


A RISC-V 32 bits Forth in 454 bytes

Alvaro G. S. Barcellos,
may 2026
github.com/agsb





*To master
riding
bicycles you
must ride
bicycles*

Proposal

A minimal functional Forth (engine and words)

Tuned for size (not for performance)

Expands the dictionary as any Forth

Fine tune Read, Eval, Print, Loop

Use of Minimal Indirect Thread Code

Current result <https://github.com/agseb/milliForth-RiscV>

454 bytes (without ELF headers) using:

88 bytes with headers, (11 Forth minimal words)

66 bytes with native syscalls, (linux)

300 bytes of native code (warp and beat)

no numbers, no names, no terminal input buffer,
no multitask, no multiuser, no rush

why do count without ELF headers ?
because barebones don't uses

System details

riscv64-linux-gnu- (*gcc, *as, *ld, *objdump, *objcopy, *readelf)

-nostartfiles -nodefaultlibs -static -Oz -march=rv32gc -mabi=ilp32


"G" is simply a shorthand for IMAFD (General Purpose)

"C" is Standard Extension for Compressed Instructions

qemu-riscv32

-display none -bios none -serial mon:stdio -smp 1 -machine virt

Could be any barebone system with _init, _getc, _putc, _exit functions



Know to tune RISCV native code

A RISCV have 32 registers of 32 bits, X0 to X31, and 32 bits opcodes.

X0 hardwired to 0x0, X1 return address (link and jump), X2 stack pointer,

Intensive use Standard Compressed Instruction Set Extension, (3 : 1)

Use X8-X15 to make 16 bits opcodes, others make 32 bits opcodes,

No advantage in use SP as Forth stack, no magic automatic push or pull,

Use a register as constant memory pointer, and use relative offsets,

Use JAL and JALR with return address register,

Move values between registers, instead use of stack to memory.

Look Ma, no terminal input buffer !

Chuck executes or compiles each word individually rather than line by line. In fact Chuck doesn't really have lines. [Jeff Fox, 2000]

No Terminal Input Buffer, only a input stream parser

A Forth engine is not an editor, no need of copy, paste, delete, insert, etc

Simply parse the tokens, valid ASCII characters between spaces,
and convert tokens into hashes.

How to be without names

WE STI.. DON.. SEE THE NEE. FOR 31 CHA..... NAM.. IN THE GEN.... CAS.

[Chuck Moore, 1983] we still don't see the need for 31 characters names in the general case

Do only hashes, no names, on linked list headers of dictionary

Link + Hash, no name length limit, small and fast,

Use 32 bits DJB2 hash, public domain, non cryptographic,

(KEY R> DUP DUP + DUP + DUP + DUP + DUP + + XOR >R)

In all words of Forth 2012 annex H, only F~ and D< have same djb2 hash



Minimal world of words

Which could be the minimum set of words? It's a recurring theme.

My choices :

1. Words that can not be composed by another words

NAND | NOR PLUS FETCH STORE NOTZEROIF

NOTZEROIF return TRUE if TOS is not zero

2. Words that must exists to compile new words

KEY EMIT COLON SEMIS EXIT USERAT

USERAT return the address of user area variables

Warp and Beat ~ 256 bytes

Engines, mostly linear native code or two level calls

main cold warm miss abort quit	\ setup	~ 56
warp (token find eval compile execute)	\ R.E.P.L.	~ 42
token (skip hash scan mask)	\ parser	~ 46
beat (unnest next pick jump nest move)	\ M.I.T.C.	~ 32
comma _init _getc _putc _exit	\ support	~ 80

Minimal vocabulary ~198 bytes

0# TOS = zero if FALSE else TRUE then	U@ pointer to SP RP LATEST DP STATE
@ fetch, load a word from memory	! store, store a word into memory
+ plus, add two values	NAND NOT AND two values
: colon, change to compile mode, etc	; semis, change to execute mode, etc
KEY read a byte from system stream	EMIT write a byte into system stream
EXIT ends a word and go to next word	

~ average of (8 + 10) bytes per word

Extra vocabulary ~ 202 bytes

This extra set can be selected during native code compilation.

NAN Not a Number, 0x80000000 to TOS	;\$ DOCODE jump and link to ROS
LSHIFT left shift 1-31 bits	RSHIFT right shift 1-31 bits
ABORT restart the outer loop warp	BYE ends the Forth
. DOT shows a hex number at TOS	\$ DOLAR receive a hex number to TOS

\$ used to parse next token as 32-bit hexadecimal value, Eg. \$ FEEDCAFE ,

More 42 primitives ~ 670 bytes

This set can be selected during native code compilation.

DUP	DROP	OVER	SWAP	ROT	LIT
AND	OR	XOR	NOT	NEG	ALIGN
R>	>R	R@	-	M*	/MOD
=	<	U<	0>	C@	C!
TRUE	FALSE	CELL	BRANCH	0BRANCH	GBRANCH
STATE	LATEST	HEAP	CEIL	FAUX	PERFORM
S0	R0	SP@	RP@	SP!	RP!

What is need for

Forth internal variables in user area (U@)

SP the pointer to user stack

RP the pointer to return stack


HEAP the address of next unused memory cell, aka DP

LAST the address of last node of dictionary linked list, aka LATEST

STAT the multi STATE of Forth, (execute, compile, [postpone, etc])

FAUX the address of HEAP when start compiling, to be next LATEST

CEIL the address of last unused memory cell



some details

Only one flag, for immediate, 0x80000000. It also is NaN for numbers.

All BRANCHs are relative. Separated stacks of 36 (?) cells deep.

The IP (s1) and USER (s0) pointers must be keep reserved.

A hash is 4 bytes, the average name length in Forth is 6 bytes.

Any word could start a compiled word. (no mandatory docol, dolit, etc)

Forth ANSI X3.215-1994, FALSE is 0x00000000, TRUE is 0xFFFFFFFF

Native multi STATE. (execute, compile, [postpone, more ?])

RAM_LIMIT = 0x40000 (64k words of 4 bytes, may be far more than this)

How deep is your Forth

	average case			worst case		
compiled 221 words	words per word	return depth	primitives per depth	word	return depth	primitives per deep
11w	6	10	277	TO	16	1942
61w	7	4	18	TO	8	72

howdeep.awk over Forth files with one word per line and same set for both cases.

any word not defined is counted as primitive

Work in Progress

Most of the slowness because variables are accessed by offsets.

Create a speed version (aeForth) with more primitives words.

Compose (from various sources) the compiled word dictionary.

Test (really) on a Raspberry Pi RP 2350 RISC-V.

Create a version for Raspberry Pi RP 2040 ARM Cortex-M0.

Test (really) on a Raspberry Pi RP 2040 ARM Cortex-M0.



That's
all
folks

How deep is your Forth

howdeep.awk is a rudimentary awk script for parsing a set of words in Forth files and counting, for each word, how many words used inside, maximum depth of stack, primitives executed, compiled executed, average primitives per stack level and listing the set of primitive types executed. Just one word per line and any word not defined is counted as primitive. Same dictionary set for both, not optimized.

	average case			worst case		
compiled 221 words	words	stack depth	primitives per depth	word	stack depth	primitives per deep
11w	6	10	277	TO	16	1942
61w	7	4	18	TO	8	72

Tuned for size not performance

What could be the minimal functional Forth engine and primitive words ?

FIGForth PDP-11 (16 bits, KA11/KD11, 1980) ~ 245 words, ? bytes

jonesForth (16 bits, X86) ~ 85 words, ? bytes, ELF

sectorForth (16 bits, x86) ~ 17 words, 512 bytes, ELF

milliForth (16bits, x86) ~13 words, 336 bytes, ELF

nybbleForth (16bits, verilog) ~12 words, ? , ?

milliForth (8 bits, 6502) ~ 11 words, 1110 bytes, barebone

milliForth (32 bits, RISCV) ~ 11 words, 454 bytes, barebone

Laziness and Complexity

Why still not KEY? in native code ?

BIOS, Basic Input Output System, 96 pages

https://www.dmtf.org/sites/default/files/standards/documents/DSP0134_2.6.1.pdf

ACPI, Advanced Configuration and Power Interface Specification, 1202 pages

https://uefi.org/sites/default/files/resources/ACPI_Spec_6.6.pdf

UEFI, Unified Extensible Firmware Interface, 2301 pages

https://uefi.org/sites/default/files/resources/UEFI_Spec_Final_2.11.pdf

RISCV, the RISC-V unprivileged architecture,

<https://docs.riscv.org/reference/isa/attachments/riscv-unprivileged.pdf>, 696 pages

no numbers and no rush.

<pre>: -1 U@ 0# ; : 0 -1 -1 NAND ; : TRUE -1 ; : FALSE 0 ; : 1 -1 -1 + -1 NAND ; : 2 1 1 + ; : 4 2 2 + ; : CELL 4 ; : SP U@ ; : RP SP CELL + ; : SP@ SP @ CELL + ; : RP@ RP @ CELL + ;</pre>	<pre>: DUP SP@ @ ; : OVER SP@ CELL + @ ; : NOT DUP NAND ; : AND NAND NOT ; : BRANCH RP@ @ DUP @ + RP@ ! ; : OBRANCH 0# NOT RP@ @ @ CELL - AND RP@ @ + CELL + RP@ ! ; : SWAP OVER OVER SP@ CELL + CELL + CELL + ! SP@ CELL + ! ; : OR NOT SWAP NOT AND NOT ; : NOR OR NOT ;</pre>	<pre><u>pull</u> LW a4, SPTR (s0); LW a3, 0 (a4); ADDI a4, a4, 4; SW a4, SPTR (s0); <u>push</u> LW a4, SPTR (s0); ADDI a4, a4, -4; SW a3, 0 (a4); SW a4, SPTR (s0);</pre>
--	--	--

Details of indirect access

Use of: LW a4, SPTR (s0); LW a3, 0 (a4); ADDI a4, a4, 4; SW a4, SPTR (s0);

Instead of: LW a3, a4; ADDI a4, a4, 4; Because can not expose registers.

Then, all variables SP, RP, LATEST, DP, STATE are relatives to user struct U@

: SP = U@ ; : SP@ SP @ ; : SP! SP ! ;

: SP+ SP@ CELL + SP! ; : SP- SP@ CELL - SP! ;

: LATEST U@ CELL + CELL + ; : LATEST@ LATEST @ ; : LATEST! LATEST ! ;

: DP U@ CELL+ CELL + CELL + ; : STATE U@ CELL + CELL + CELL + CELL + ;

Jump and Link (32 bits)

Jump to subroutine

; push RA, one deep

$SP \leftarrow SP - 4$

$[SP] \leftarrow RA$

; link and jump

$RA \leftarrow PC + 4$

$PC \leftarrow EA$

Return from subroutine

; pull RA, one deep

$RA \leftarrow [SP]$

$SP \leftarrow SP + 4$

; link and jump

$ZR \leftarrow PC$

$PC \leftarrow RA$

PC, program counter; EA, effective code address; the registers defaults to ZR (x0), wired zero; RA (x1), return address; SP (x2), stack pointer;

References

<https://groups.google.com/g/comp.lang.forth/c/NS2icrCj1jQ>

<https://archive.org/details/the-road-towards-a-minimal-forth-architecture>

https://www.forth.org/fig-forth/fig-forth_PDP-11.pdf

<https://rwmj.wordpress.com/2010/08/07/jonesforth-git-repository/>

<https://github.com/cesarblum/sectorforth>

<https://github.com/fuzzballcat/milliForth>

<https://github.com/agsb/milliForth-6502>

<https://github.com/agsb/milliForth-RiscV>

References

<https://en.wikipedia.org/wiki/RISC-V>

<https://theforth.net/package/minimal>

<https://github.com/larsbrinkhoff/nybbleForth>


<https://github.com/meithecatte/miniforth/tree/post1>

<https://github.com/larsbrinkhoff/nybbleForth>

<https://www.youtube.com/watch?v=Ok-1E8xwo90>

<https://www.greenarraychips.com/home/documents/greg/DB001-171107-F18A.pdf>






#D14734

#E78570

#CCD6FB

#6C8FB5




#614947

#B6685B

#FDA369

#FDD6B6




#AD8997

#5B3A57

#E95223

#FC952D



#473B06

#DD9F0D

#7DD6EA

#E1F8FF



#E68F59

#073E6B

#1973AE

#A1D8EC




#333C31

#97B45D

#A49752

#D7B56B



#0189DD

#04BCF2

#8AD9F7

#D2B178




#F11532

#F046A1

#FA615B

#F97D03




#4F566E

#8C9EB4

#F79B7E

#FFECA7




#7B574B

#D3855F

#FEBF72

#FFF0B6




#3C1C0F

#855615

#AD972F

#3B4D0C



#36352F

#465321

#B59226

#DF560E